SANDEEL FARMS | Minimum-Cost OpenAI Plan

Developer handoff • 10 October 2026 • Pricing in USD; API charges separate from ChatGPT subscriptions.

Recommended starter configuration

Text model: GPT-6 Luna (**gpt-6-luna**) for all four functions. **Voice input:** GPT-4o Mini Transcribe (**gpt-4o-mini-transcribe**) when needed. **Voice output:** free browser text-to-speech where supported. **Realtime voice:** OFF initially.

Why: This provides a low-priced, current general assistant without building expensive continuous audio sessions. GPT-5 nano is cheaper on paper but is marked deprecated, so avoid it for a new deployment.

1 | Text AI pricing (standard processing)

Model	Input / 1M tokens	Output / 1M tokens	Decision
GPT-6 Luna	$0.10	$0.50	RECOMMENDED
GPT-5.6 Luna	$0.20	$1.20	Alternative
GPT-4o Mini	$0.15	$0.60	Alternative
GPT-5 nano	$0.05	$0.40	Deprecated
GPT-5.6 Sol	$4.00	$20.00	Not needed

2 | Voice options and costs

Voice service	Published rate	Decision
GPT-4o Mini Transcribe	$0.003 / audio minute	Lowest quoted transcription
GPT-Transcribe	$0.0045 / audio minute	Alternative; test Urdu quality
GPT-4o Mini TTS	$0.60 / 1M text input tokens + $12 / 1M audio output tokens	Optional; deprecated
Realtime 2.1 Mini	$10 / 1M audio input tokens + $20 / 1M audio output tokens	Keep OFF initially

3 | Four AI functions to add to Sandeel Farms

1. AI Farm Assistant — Ask for herd counts, breeding eligibility, milk totals, feed stock and expense balances. Read authorized live farm data; do not invent records.

2. Voice Assistant — Speak commands in English, Urdu or Roman Urdu. Transcribe audio, run the same approved assistant tools, return text and optionally browser-spoken answers.

3. AI Record Management — Draft goat updates, milk entries, feed transactions and breeding changes. Check roles, request confirmation for protected writes, validate and audit all changes.

4. AI Farm Reports — Summarize breeding groups, pregnancy outcomes, milk yields, feed use and farm finances using actual database figures; support date filters.

4 | Example monthly API budget

Assumptions: 1,000 assistant questions/month; average 1,000 input + 300 output tokens/question; 100 minutes of transcription. GPT-6 Luna text = **$0.25**; GPT-4o Mini Transcribe = **$0.30**; **estimated total = $0.55/month**. Using GPT-Transcribe instead = **$0.70/month**. Actual prompts, history, tool calls, retries and output can raise costs. Browser speech is not guaranteed on every device.

5 | Developer instructions and controls

Use the OpenAI Responses API from the secure Node backend only; never place an API key in the HTML. Keep existing roles, tenant isolation, approval prompts, and audit logs. Start with a $5/month usage alert and a backend spending guard. Disable optional web search, hosted tools, realtime audio and paid TTS by default. Log token and transcription usage per user; add admin-configurable model selection. Test Urdu, Roman Urdu and mixed-language commands before launch.

Official sources (checked 10 October 2026)

Text: https://developers.openai.com/api/docs/pricing | https://developers.openai.com/api/docs/models/gpt-6-luna
Transcription: https://developers.openai.com/api/docs/models/gpt-4o-mini-transcribe | https://developers.openai.com/api/docs/models/gpt-transcribe
TTS: https://developers.openai.com/api/docs/models/gpt-4o-mini-tts

Notes: Pricing may change; verify before deployment. No live API billing or working deployment has been confirmed by this document.